

PROOFS/NOTES

NIKHIL VASAN

1. OVERVIEW

2. IMPLEMENTATION

The implementation of the burning of the baseFee will be done through introducing a dependency between the Distribution Keeper, and feeMarket Keeper as such, in x/Distribution/types/interfaces.go

```
//code example
type FeeMarketKeeper interface{
    // fields entered here ..
}
```

3. GFS

Architecture. Master is defined as follows

```
type Master {
    // chunkData stored in memory
    chunkHandles map[FID]*chunkData
}
// identifier for files stored in GFS
type FID = string
// definition of chunkData, maintains the data required for
// master when granting
// access to chunkServer for chunk requested
type chunkData {
    // essentially IP address of chunk server
    servers []*chunkServer
    // primary, server with current lease for sequencing record
    appends / writes
    primary *chunkServer
}
```

GFS handles writes / reads to chunks as follows.

- (1) Client sends request to chunk server, with FID / offset
- (2) Client identifies chunk for offset, sends primary for lease to client
- (3) Client caches primary address, and transmits data as well as operation to be sequenced by primary

In writes / appends, the primary is tasked with sequencing the operation to the file for all replicas. In each case, it must be required that the data is atomically appended (not necessarily at end of chunk), to file, primary sequences concurrent writes, and transmits sequence to replicas. All replicas gossip data to each replica in sequence ($rep1 \Rightarrow rep2 \Rightarrow \dots$), where rep_i is an ordering determined by latency, etc.

The Master maintains a log of all writes / reads. Readers of chunks obtain, read locks, ensure that version number for file is stored. Files with read-lock, ensure that name-space changes / leases to primary for writes cannot be executed until read-locks released. How to handle?

4. DISTRIBUTED ALGORITHMS

- fail-stop* - Any process can fail, but failures can be reliably detected by other processes.
 - fail-silent* - Process crashes never can be detected
 - fail-noisy* - Identification of Process failure is only eventual
 - fail-recovery* - processes can fail / recover and participate in alg.
 - fail-arbitrary* - processes can fail in perhaps arbitrary ways
 - randomized* - Processes may make probabilistic choices given a common source of randomness
- Denote set of processes in algorithm as Π , define *rank*, a mapping $rank : \Pi \rightarrow \{1, \dots, N\}$
Processes send messages denoted $m_{p,i}$, where i is a *sequence number*, notice, combined w/ rank, a total ordering is defined (eg. *lexicographic ordering on rank, sequence num.*) over the messages